

# Lecture 6 – Of Trees and Forests

Martin Spindler

2016-03-15

# Introduction

- ▶ Idea: Partition the feature / covariables space into a set of rectangles and then fit a simple model (constant) in each one. So the estimated function is the average of outcomes falling in this rectangle.
- ▶ Recursive binary partitions, i.e. sequentially we choose variable and corresponding split point which achieve best fit until some stopping criterion is reached.
- ▶ Here: Regression, but also used for classification (with different criteria)
- ▶ Example [cf blackboard]

# Regression Trees

$n$  observations  $(x_i, y_i), i = 1, \dots, n, \quad x_i = (x_{i1}, \dots, x_{ip})$

Given a partition into  $M$  regions  $R_1, \dots, R_M$  with a fitted constant in each region:

$$f(x) = \sum_{m=1}^M c_m I(x \in R_m)$$

Minimizing the residual sum of squares (RSS)  $\sum_{x_i \in R_m} (y_i - f(x_i))^2$  leads to  $\hat{c}_m = \text{ave}(y_i | x_i \in R_m)$

Finding the best partition in terms of minimal RSS is generally computational infeasible.

Instead: greedy algorithm

# Regression Trees

## Algorithm:

Start with all data

1. We consider a splitting variable  $j$  and split point  $s$ , s.t.

$$\min_{j,s} \left[ \min_{c_1} \sum_{x_i \in R_1(j,s)} (y_i - c_1)^2 + \min_{c_2} \sum_{x_i \in R_2(j,s)} (y_i - c_2)^2 \right]$$

with  $R_1(j, s) = \{X|X_j \leq s\}$ ,  $R_2(j, s) = \{X|X_j > s\}$

Determination of the best pair  $(j, s)$  is feasible.

2. Repeat 1) on each of the two resulting regions
3. Stop when some criterion is reached

# Regression Trees

How large should we grow a tree?

Tree size is a tuning parameter governing the model's complexity. It should be chosen adaptively from the data.

Strategy 1: Split tree node only, if decrease in rss is sufficiently high (but too short-sighted)

Strategy 2: (preferred)

- ▶ Given a large tree, stopping when some minimal node size is reached
- ▶ Prune the tree by *cost-complexity pruning*

# Regression Trees

Subtree  $T \subset T_0$  is any tree that can be obtained by pruning  $T_0$ , i.e. collapsing any number of its internal nodes. We denote the terminal nodes by  $m = 1, \dots, M$ .

$$N_m = |\{x_i \in R_m\}|$$

$$\hat{c}_m = 1/N_m \sum_{x_i \in R_m} y_i$$

$$Q_m(T) = \sum_{x_i \in R_m} (y_i - \hat{c}_m)^2$$

$$\text{Cost-complexity criterion } C_\alpha(T) = \sum_{m=1}^M N_m Q_m(T) + \alpha |T|$$

# Regression Trees

Find  $T_\alpha \subset T_0$  to minimize  $C_\alpha(T)$

$\alpha$  governs the trade-off between tree-size and its goodness of fit to the data.

$\alpha = 0$  yields full tree.

Choice of  $\alpha$  by cross validation.

Final tree  $T_{\hat{\alpha}}$

# Bagging

- ▶ Bagging: Bootstrap aggregation or bagging averages
- ▶ Training data  $Z = \{(x_1, y_1), \dots, (x_n, y_n)\}$
- ▶ Fit a model to  $Z$  and obtain  $\hat{f}(x)$
- ▶ Idea: average the predictions over a collection of bootstrap samples, thereby reducing its variance
- ▶
- ▶  $Z^{*b}, b = 1, \dots, B$  bootstrap samples
- ▶ Fit model to get  $\hat{f}^{*,b}(x)$
- ▶  $\hat{f}_{bag}(x) = 1/B \sum_{b=1}^B \hat{f}^{*,b}(x)$



# Bagging

- ▶  $\hat{f}_{bag}(x)$  estimate of the true bagging value  $E_{\hat{p}}\hat{f}^*(x)$
- ▶ Well suited for high-variance, low-bias procedures
- ▶ Application: regression trees

# Random Forests

- ▶ Introduced by Breiman (2001)
- ▶ Very powerful (good performance) in many applications
- ▶ Modified version of bagging
- ▶ Idea: Building a large collection of de-correlated trees and then average them (Breiman, 2001)

# Random Forests

- ▶ Trees: low bias, but very noisy / high variance  $\Rightarrow$  goal : reduction of variance
- ▶ Trees generated by bagging are identically distributed, but not necessarily independent
- ▶ For identical distributed variables with positive pairwise correlation  $\rho$ : variance of average  $\rho\sigma^2 + (1 - \rho)/B\sigma^2$  ( $\rho$  correlation of the trees) (for i.i.d. rvs:  $\sigma^2/B$ )
- ▶ Application: nonlinear estimators like random trees

## Random Forests | Procedure

- ▶ Bootstrap samples  $1, \dots, B$
- ▶ Build trees and  
“Before each split, select  $m \leq p$  of the input variables at random as candidates for splitting” (e.g.  $m = \sqrt{p}$ ,  $m = 1$ )
- ▶ Aggregation:

$$\hat{f}_{rf}(x) = \frac{1}{B} \sum_{b=1}^B T(x; \Theta_b)$$

$\Theta_b$ : split variables, cut points, terminal node values for  $b$

# Ensemble Learning | Introduction

- ▶ Idea: To build a prediction model by combining the strengths of a collection of simpler base models.
- ▶ Bagging, random forests are ensemble methods for classification, where a committee of trees each cast a vote for predicted class.
- ▶ Boosting proposed as a committee method where the committee of weak learners evolves over time with a weighted vote for the members
- ▶ Ensemble learning
- ▶ Developing a population of base learners from the training data
- ▶ Combining them to form the composite predictor

# Learning Ensembles

- ▶ We consider functions of the form

$$f(x) = \alpha_0 + \sum_{T_k \in \mathcal{T}} \alpha_k T_k(x)$$

with  $\mathcal{T}$  dictionary of basis functions, e.g. trees with  $|\mathcal{T}|$  quite large.

- ▶ Hybrid approach of Friedman and Popescu (2003)
- ▶ A finite dictionary  $\mathcal{T}_L = \{T_1(x), \dots, T_M(x)\}$  of basis functions is induced from the training data.
- ▶ A family of functions  $f_\lambda(x)$  is built by fitting a lasso path in this dictionary

$$\alpha(\lambda) = \underset{\alpha}{\operatorname{argmin}} \sum_{i=1}^N L[y_i, \alpha_0 + \sum_{m=1}^M \alpha_m T_m(x_i)] + \lambda \sum_{m=1}^M |\alpha_m|$$

# Ensemble Generating Algorithm

How to choose the set of base functions  $b(x; \gamma)$  forming  $\mathcal{T}_L$ ?

- ▶  $f_0(x) = \arg \min_c \sum_{i=1}^N L(y_i, c)$
- ▶ For  $m = 1$  to  $M$  do
- ▶  $\gamma_m = \arg \min_{\gamma} \sum_{i \in S_m(\eta)} L(y_i, f_{m-1}(x_i) + b(x_i; \gamma))$
- ▶  $f_m(x) = f_{m-1}(x) + \nu b(x; \gamma_m)$
- ▶  $\mathcal{T}_{ISLE} = \{b(x; \gamma_1), \dots, b(x; \gamma_M)\}$

$S_m(\eta)$  refers to a subsample of  $N\eta$  of the training observations, typically without replacement.

Recommendation:  $\eta \leq 1/2$  and  $\eta \sim 1/\sqrt{(N)}$ ,  $\nu = 0.1$

*Importance sampled learning ensemble (ISLE)*

## Special cases of the Algorithm

- ▶ Bagging:  $\eta = 1$ , samples with replacement,  $\nu = 0$
- ▶ Random forest: sampling is similar with more randomness introduced by the selection of the splitting variable.
- ▶ Gradient boosting with shrinkage uses  $\eta = 1$
- ▶ Stochastic gradient boosting: identical