

Lecture 9 – Support Vector Machines

Martin Spindler

2016-04-22

Introduction

- ▶ Outcome variable: $Y \in \{-1, +1\}$
- ▶ Generalization to multi-valued outcomes is straight forward.
- ▶ Basic Idea (for classification): Separating the space of features by hyperplanes into different regressions so that the dependent variable / outcome is separated.
- ▶ *Seperating* hyperplanes
- ▶ Graph: cf white board
- ▶ In this lecture we focus only on the basic idea.

Digression: Hyperplanes

- ▶ In a p -dimensional space a *hyperplane* is a flat affine subspace of dimension $p-1$.
- ▶ Example 1: In $p = 2$ a hyperplane is a line.
- ▶ Example 2: In $p = 3$ a hyperplane is a plane.
- ▶ A hyperplane in a p -dimensional space is defined by

$$\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p = 0$$

- ▶ If $\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p < 0$ then $X = (X_1, \dots, X_p)$ lies on one side of the hyperplane.
- ▶ If $\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p > 0$ then X lies on the other side of the hyperplane.
- ▶ Hence: Separation of the space into two parts.
- ▶ $f(x) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$ gives the signed distance from a point x to the hyperplane defined by $f(x) = 0$

Support Vector Classifier

- ▶ Training data: $(x_1, y_1), \dots, (x_n, y_n)$, $x_i \in \mathbb{R}^p$ and $y_i \in \{-1, +1\}$
- ▶ Definition hyperplane: $\{x : f(x) = x^T \beta + \beta_0 = 0\}$ (β unit vector with $\|\beta\| = 1$)
- ▶ Classification rule: $G(x) = \text{sign}[x^T \beta + \beta_0]$

Support Vector Classifier: Separable Case

- ▶ Separable means: $y_i f(x_i) > 0 \forall i$ for a plane $f(x)$
- ▶ Find the hyperplane that creates the biggest margin between the classes for -1 and $+1$.
- ▶ Optimization problem: $\max_{\beta, \beta_0, \|\beta\|=1} M$ subject to

$$y_i(x_i^T \beta + \beta_0) \geq M, i = 1, \dots, n$$

- ▶ Equivalent formulation: $\min_{\beta, \beta_0} \|\beta\|$ subject to

$$y_i(x_i^T \beta + \beta_0) \geq 1, i = 1, \dots, n$$

Support Vector Classifiers: Non-separable Case

- ▶ Now: classes overlap in the feature space.
- ▶ Still maximize M , but allow for some points to be on the wrong side of the margin.
- ▶ Slack variables $\xi = (\xi_1, \dots, \xi_n)$
- ▶ Modification of the constraints: $y_i(x_i^T \beta + \beta_0) \geq M - \xi_i$ or $y_i(x_i^T \beta + \beta_0) \geq M(1 - \xi_i)$
- ▶ $\forall \xi \geq 0, \sum_{i=1}^n \xi_i \leq \text{constant}$
- ▶ Interpretation: overlap in actual distance from the margin vs overlap in relative distance.
- ▶ Focus on the second case (b/c convex optimization problem)

Support Vector Classifiers: Non-separable Case

- ▶ ξ_i proportional amount by which the prediction $f(x_i)$ is on the wrong side of the margin.
- ▶ Missclassification occurs, if $\xi_i > 1$
- ▶ Bounding $\sum \xi_i$ at value K , bounds the number of training missclassifications at K .
- ▶ Equivalent formulation of the problem: $\min \|\beta\|$ subject to

$$y_i(x_i^T \beta + \beta_0) \geq 1 - \xi_i \forall i, \xi_i \geq 0, \sum \xi_i \leq K$$

Support Vector Machines

- ▶ Up to now: linear boundaries in the feature space.
- ▶ Flexibility by enlarging the feature space using basis expansions (e.g. polynomials, splines)
- ▶ Better training-class separation and nonlinear boundaries in the original space.
- ▶ Selection of basis functions $h_m(x)$, $m = 1, \dots, M$ and fit of SV classifier using the input features $h(x_i) = (h_1(x_i), \dots, h_M(x_i))$.
- ▶ Nonlinear function $\hat{f}(x) = h(x)^T \hat{\beta} + \hat{\beta}_0$
- ▶ Classifier: $\hat{G} = \text{sign}(\hat{f}(x))$

Support Vector Machines

- ▶ SVM use a very large space of basis functions leading to computational problems.
- ▶ Problem of overfitting.
- ▶ SVM technology takes care of both problems.

Support Vector Machines

- ▶ (omitting technical details)
- ▶ Solution of the optimization problem involve $h(x)$ only through inner products.
- ▶ Knowledge of the kernel functions $K(x, x') = \langle h(x), h(x') \rangle$ is sufficient.
- ▶ Examples: + d th degree polynomial: $K(x, x') = (1 + \langle x, x' \rangle)^d$
+ Radial basis: $K(x, x') = \exp(-\gamma \|x - x'\|^2)$

Support Vector Machines | Example

- ▶ Consider two-dimensional space (X_1, X_2) and polynomial kernel of degree 2.
- ▶ $K(X, X') =$
 $1 + 2X_1X'_1 + 2X_2X'_2 + (X_1X'_1)^2 + (X_2X'_2)^2 + 2X_1X'_1X_2X'_2$
- ▶ Then $M = 6$ and $h_1(X) = 1, h_2(X) = \sqrt{2}X_1, \dots$. Then
 $K(X, X') = \langle h(X), h(X') \rangle$

SVM as a Penalization Method

- ▶ With $f(x) = h(x)^T \beta + \beta_0$, we consider the optimization problem:

$$\min_{\beta_0, \beta} \sum_{i=1}^n [1 - y_i f(x_i)]_+ + \frac{\lambda}{2} \|\beta\|^2$$

- ▶ loss + penalty
- ▶ Hinge loss function: $L(y, f) = [1 - yf]_+$
- ▶ Solution to the above optimization problem (with $\lambda = 1/C$) is the same as for the SVM problem.
- ▶ (C is a Cost parameter related to K)

SVM | An Illustration

We use the library *e1071* (alternative *LiblineaR* for very large linear problems)

```
library(e1071)
```


```
## Warning: package 'e1071' was built under R version 3.1.3
```

First, we generate data we would like to classify

```
set.seed(12345)
x = matrix(rnorm(200*2), ncol=2)
x[1:100,] = x[1:100,] + 2
x[101:150,] = x[101:150,] - 2
y = c(rep(1,150), rep(2,50))
dat = data.frame(x=x, y=as.factor(y))
```

SVM | An Illustration

```
plot(x, col=y)
```



SVM | An Illustration

Next, we split the data into training and testing sample


```
train = sample(200,100)
```

Then we fit a SVM with radial basis and plot the result

```
svmfit = svm(y~., data=dat[train,], kernel="radial", gamma=  
# summary(smvfit)
```

SVM | An Illustration

```
plot(svmfit, dat[train,])
```



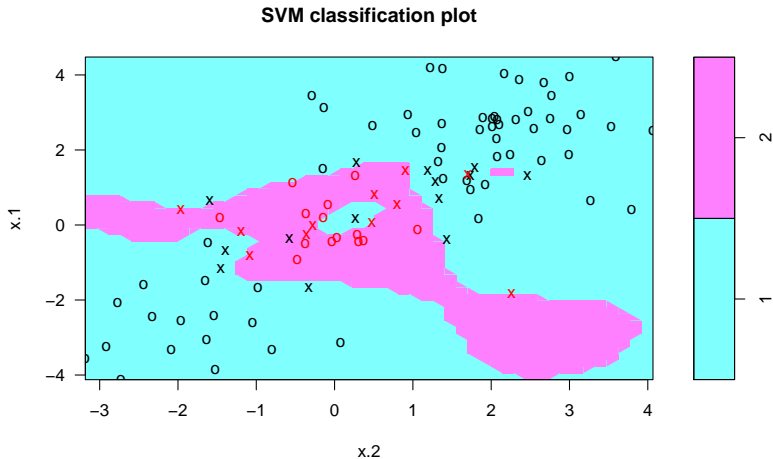
SVM | An Illustration

We can now increase the cost parameter to reduce the training errors

```
svmfit = svm(y~., data=dat[train,], kernel="radial", gamma=  
#plot(svmfit, dat[train,])  
# summary(svmfit)
```

SVM | An Illustration

```
plot(svmfit, dat[train,])
```



SVM | An Illustration

Selection of the cost parameter and γ by CV

```
tune.out=tune(svm, y~., data=dat[train,], kernel="radial",  
#summary(tune.out)
```

SVM | An Illustration

Finally, we test it on the testing data

```
table(true=dat[-train,"y"], pred=predict(tune.out$best.model
```

```
##      pred
## true  1  2
##      1 70  5
##      2  4 21
```